

## Lab 5 Report

Output:

-----  
Image Info ::

Height=246 Width=300

Time taken by master process during convergence loop: 1.49 s

Sum of time taken by all MPI processes during convergence loop: 23.89 s

Threshold: 49  
-----

<i>coins.bmp</i>	Serial	Cuda	MPI (and OMP)
<b>Time</b>	100 ms	3.7 ms	1.49 s

Other MPI results:

imgs	Time (taken by master process) (s)	Time taken by all processes (s)
fox.bmp	1.43	22.90
map.bmp	2.79	44.74
cave.bmp	2.93	46.96
cliff.bmp	3.77	60.44

- **Compare your run-time with your serial and cuda parallel versions from lab4. Did you achieve the best performance with the serial or a parallel version? Why?**

For coins.bmp

The Cuda parallel version gave the best run-time (3.7 ms) followed by serial version (100 ms) followed by MPI (1.49 s). MPI was slower than Cuda as all the work was done on only 1 GPU. Overall, MPI gave the worst run-time amongst the three and the advantage of using it was not apparent for such a small file as coins.bmp.

- **Explain your workload distribution strategy and why you selected it.**

The program was designed using block distribution. Each MPI process handled (rank \* quotient + 1) 'i' values from serial version (where quotient = (height – 2)/# of processes).

The inner j-for loop was parallelized using 2 (which gave the best results; see table below) OMP threads for each process.

This implementation provides the advantage of not using any barriers in program which significantly improves the running time.

on cliff.bmp (83.18 MB)

# OMP threads	Time (taken by master process) (s)	Time taken by all processes (s)
2	3.77	60.44
4	5.86	93.87
8	29.36	469.88
16	71.68	1147.55

32	5.03	80.47
64	8.39	134.19
128	16.19	259.03
256	32.82	525.09

- **Report and comment on any reduction and/or synchronization mechanisms you used.**

The program reported aggregate time taken by all MPI processes calculated using MPI\_Reduce. MPI\_Reduce and OMP reduction were used during convergence loop to sum black\_cells into percent\_black\_cells which was then broadcasted to all processes using MPI\_Broadcast to check the condition of convergence loop.

At the end of convergence loop, all the slave processes sent their corresponding new pixel values to master process using MPI\_Send who collected them using MPI\_Recv.

- **Were there any surprises you encountered in this exercise?**

No.

Note: The extra test images are available here:- <https://osu.app.box.com/v/CSE5441TestImages>